

RAG Knowledge Base - Page 1

RAG stands for Retrieval Augmented Generation.

A RAG chatbot first searches relevant documents, then uses the retrieved context to answer.

The quality of the answer depends on document loading, chunking, embedding, and retrieval.

Document Loading - Page 2

Document loaders convert files into Document objects.

Each Document contains `page_content` and `metadata`.

Metadata such as `source` and `page` helps users verify the answer source.

Text Splitting - Page 3

Text splitters divide long documents into searchable chunks.

Small chunks may lose context, while large chunks may include noise.

`RecursiveCharacterTextSplitter` is a good starting point for general PDF RAG.

Embedding - Page 4

Embedding models convert text into vectors.

Similar meanings become close vectors.

Vector stores use these vectors for semantic search.